

Consistency and classification of metrics for binary classifiers

K. Dirland
<***@***>

A. S. Lundervold
<***@***>
(or any permutation thereof)

P.G.L. Porta Mana 
<pgl@portamana.org>

Draft. 4 March 2022; updated 8 April 2022

abstract

✚ [Luca] I find it very difficult to structure the paper: there seems to be issues at several levels in the development and use of binary classifiers (and classifiers in general) within machine-learning. Here are some relevant points:

- There should be a distinction between ‘inference’ (or forecast, prediction, guess) and ‘decision’ (or action, choice). In particular, the possible situations we may be uncertain about and the possible decisions available may be completely different things. A clinician, for example, may be uncertain about ‘cancer’ vs ‘non-cancer’, while the choices are about ‘drug treatment 1’ vs ‘drug treatment 2’ vs ‘surgery’.
- Probability theory & decision theory say that in order to make self-consistent decision we need two things: (a) the probabilities for the possible situations, (b) the utilities of the decisions given each possible situation.
- A useful machine-learning algorithm should therefore give us one of two things:
 - either the *probabilities* of the uncertain situations (‘cancer’ vs ‘non-cancer’ in the example above),
 - or the final decision (‘drug treatment 1’ vs ‘drug treatment 2’ vs ‘surgery’ in the example above).

Current machine-learning classifiers do not give us either: the output in the example above would be ‘cancer’ vs ‘non-cancer’, often without probabilities.

- So there are two possible solutions to the problem above:
 - We must build a classifier that outputs probabilities. The 0–1 outputs of current classifiers cannot properly interpreted as probabilities, for various reasons.
 - We must build a classifier that output *decisions*: so not ‘cancer’ vs ‘non-cancer’, but ‘drug treatment 1’ vs etc..

1 Valuation metrics, amounts of data, inferences, and decisions

Let’s consider the simple example of a binary classifier and several dilemmas that appear in its development, choice, and use.

At the moment of evaluating different classifier algorithms, or different hyperparameter settings for one algorithm, we are avalanched by a choice of possible evaluation scales: accuracy, area under curve,

F_1 -measure, mean square contingency¹ also known as Matthews correlation coefficient², precision, recall, sensitivity, specificity, and many others³. Only vague guidelines are usually given to face this choice. A thorough analysis and discussion of several such scales was given by Goodman & Kruskal (1954; 1959; 1963; 1972).

We can also ask: are all these scales well-founded and self-consistent? Is it possible that the use of any of them leads to contradictions? The literature abounds with studies showing that some scale X may imply hidden contradictions with the data or the assumptions used for our inference, and is therefore worse than some other scale Y . See for example Baker & Pinsky (2001), Lobo et al. (2008), Hand & Christen (2018), Zhu (2020) and Goodman & Kruskal's papers cited above for instances of criticisms of area under the curve, F_1 -measure, Matthews correlation coefficient, and other scales.

If we have many more data for one class than for the other – a common predicament in medical applications – we must face the 'class-imbalance problem': the classifier ends up classifying all data as belonging to the more numerous class⁴, which may be an undesirable action if the misclassification of cases belonging to the less numerous class entails high costs.

The three points above turn out to be tightly related and to have a common solution. We show that

1. the admissible valuation scales for a binary classifier form a two-dimensional family; that is, the choice of a specific scale corresponds to the choice of two numbers. Such choice is problem-dependent and cannot be given a priori.
2. admissible scales are only those that can be expressed as a linear function of the elements of the population-normalized confusion matrix. Scales such as the F_1 -measure or the Matthews correlation coefficient are therefore inad***

2 Overview of decision theory

5

Decision theory makes a distinction between

¹ Yule 1912 denoted ' r ' there. ² Matthews 1975. ³ Sammut & Webb 2017. ⁴ Sammut & Webb 2017; Provost 2000. ⁵ Russell & Norvig 2022 ch. 15; Jeffrey 1965; North 1968.

- a. the possible situations we are uncertain about
- b. the possible choices we can make.

This distinction is important, in fact in some cases the numbers of possible uncertain situations

Appendix: broader overview of binary classification

Let us consider our binary-classification problem from a general perspective and summarize how it would be approached and solved from first principles⁶ if our computational resources had no constraints.

In our long-term task we will receive ‘units’ of a specific kind; the units for example could be gadgets, individuals, or investment portfolios. Each new unit will belong to one of two classes, which we can denote $X=0$ and $X=1$; for example they could be ‘defective’ vs ‘non-defective’, ‘ill’ vs ‘healthy’. The class will be unknown to us. For each new unit we shall need to decide among two possible actions, which we can denote $A=\hat{0}$ and $A=\hat{1}$; for example ‘discard’ vs ‘keep’, or ‘treat’ vs ‘dismiss’. The utility of each action depends on the unknown class of the unit; we denote these utilities by $U(A | X)$. For each new unit we will be able to measure a ‘feature’ Z of a specific kind common to all units; for example Z could be a set of categorical and real quantities, or an image such as a brain scan. We have a set of units – our ‘sample units’ or ‘sample data’ – that are somehow “representative” of the units we will receive in our long-term task⁷. we know both the class and the feature of each of these sample units. Let us denote this sample information by D .

According to the principles of decision theory and probability theory, for each new unit we would proceed as follows:

1. Assign probabilities to the two possible values of the unit’s class, given the value of the unit’s feature $Z=z$, our sample data D , and any other available information:

$$p(X=0|Z=z, D), \quad p(X=1|Z=z, D) \equiv 1-p(X=0|Z=z, D), \quad (1)$$

according to the rules of the probability calculus.

⁶ Russell & Norvig 2022 part IV. ⁷ for a critical analysis of the sometimes hollow term ‘representative’ see Kruskal & Mosteller 1979a,b,c; 1980.

2. Calculate the expected utilities \bar{U} of the two possible actions:

$$\begin{aligned}\bar{U}(\hat{0}) &:= U(\hat{0} \mid X=0) p(X=0 \mid Z=z, D) + U(\hat{0} \mid X=1) p(X=1 \mid Z=z, D) \\ \bar{U}(\hat{1}) &:= U(\hat{1} \mid X=0) p(X=0 \mid Z=z, D) + U(\hat{1} \mid X=1) p(X=1 \mid Z=z, D)\end{aligned}\tag{2}$$

and choose the action having maximal expected utility.

How is the probability $p(X \mid Z=z, D)$ determined by the probability calculus? Here is a simplified, intuitive picture. First consider the case where the feature Z can only assume a small number of possible values, so that many units can in principle have the same value of Z .

Consider the collection of all units having $Z = z$ that we received in the past and will receive in the future. Among them, a proportion $F(X=0 \mid Z=z)$ belong to class 0, and a proportion $1 - F(X=0 \mid Z=z) \equiv F(X=1 \mid Z=z)$ to class 1. For example these two proportions could be 74% and 26%. Our present unit with $Z=z$ is a member of this collection. The probability $p(X=0 \mid Z=z)$ that our unit belongs to class 0, given that its feature has value z , is then intuitively equal to the proportion $F(X=0 \mid Z=z)$. Analogously for $X=1$.

The problem is that we do not know the proportion $F(X=0 \mid Z=z)$. However, we expect it to be roughly equal to the analogous proportion seen in our sample data; let us denote the latter by $F_s(X=0 \mid Z=z)$:

$$F(X=0 \mid Z=z) \sim F_s(X=0 \mid Z=z) .\tag{3}$$

this is indeed what we mean by saying that our sample data are ‘representative’ of the future units. Later we shall discuss the case in which such representativeness is of different kinds. We expect the discrepancy between $F(X=0 \mid Z=z)$ and $F_s(X=0 \mid Z=z)$ to be smaller, the larger the number of sample data. Vice versa we expect it to be larger, the smaller the number of sample data.

If Z can assume a continuum of values, as is the case for a brain scan for example, then the collection of units having $Z=z$ is more difficult to imagine. In this case each unit will be unique in its feature value – no two brains are exactly alike.

old text below

Given the unit’s feature Z we will assign probabilities to the possible values of the unit’s class: according to the rules of the probability calculus.

As mentioned in § 2, a decision problem under uncertainty is conceptually divided into two steps

The Suppose we have a population of units or individuals characterized by a possibly multidimensional variable Z and a binary variable $X \in \{0, 1\}$. Different joint combinations of (X, Z) values can appear in this population. Denote by $F(X=x, Z=z)$, or more simply $F(x, z)$ when there is no confusion, the number of individuals having specific joint values $(X=x, Z=z)$. This is the absolute frequency of the values (x, z) . We can also count the number of individuals having a specific value of $Z=z$, regardless of X ; this is the marginal absolute frequency $F(z)$. It is easy to see that

$$F(z) = F(X=0, z) + F(X=1, z) \equiv \sum_x F(x, z). \quad (4)$$

Analogously for $F(x)$.

Select only the subpopulation of individuals that have a specific value $Z=z$. In this subpopulation, the *proportion* of individuals having a specific value $X=x$ is $f(x | Z=z)$. This is the conditional relative frequency of x given that z . It is easy to see that

$$f(x | z) = \frac{F(x, z)}{F(z)}. \quad (5)$$

Now suppose that we know all these statistics about this population. An individual coming from this population is presented to us. We measure its Z and obtain the value z . What could be the value of X for this individual? We know that among all individuals having $Z=z$ (and the individual before us is one of them) a proportion $f(x | z)$ has $X=x$. Thus we can say that there is a probability $f(x | z)$ that our individual has $X=x$. And this is all we can say if we only know Z .

For this individual we must choose among two actions $\{a, b\}$. The utility of performing action a if the individual has $X=x$, and given any other known circumstances, is $U(a | x)$; similarly for b . If we knew the value of X , say $X=0$, we would simply choose the action leading to maximal utility:

$$\begin{aligned} \text{if } U(a | X=0) > U(b | X=0) & \text{ then choose action } a, \\ \text{if } U(a | X=0) < U(b | X=0) & \text{ then choose action } b, \\ \text{else} & \text{ it does not matter which action is chosen.} \end{aligned} \quad (6)$$

But we do not know the actual value of X . We have probabilities for the possible values of X given that $Z=z$ for our individual. Since X is uncertain, the final utilities of the two actions are also uncertain; but we can calculate their *expected* values $\bar{U}(a \mid Z=z)$ and $\bar{U}(b \mid Z=z)$:

$$\begin{aligned}\bar{U}(a \mid z) &:= U(a \mid X=0) f(X=0 \mid z) + U(a \mid X=1) f(X=1 \mid z) , \\ \bar{U}(b \mid z) &:= U(b \mid X=0) f(X=0 \mid z) + U(b \mid X=1) f(X=1 \mid z) .\end{aligned}\quad (7)$$

Decision theory shows that the optimal action is the one having the maximal expected utility. Our choice therefore proceeds as follows:

$$\begin{aligned}\text{if } \bar{U}(a \mid z) &> \bar{U}(b \mid z) \text{ then choose action } a, \\ \text{if } \bar{U}(a \mid z) &< \bar{U}(b \mid z) \text{ then choose action } b, \\ \text{else } &\text{it does not matter which action is chosen.}\end{aligned}\quad (8)$$

The decision procedure just discussed is very simple and does not need any machine-learning algorithms. It could be implemented in a simple algorithm that takes as input the full statistics $F(X, Z)$ of the population, the utilities, and yields an output according to (8).

Our main problem is that the full statistics $F(X, Z)$ is almost universally not known. Typically we only have the statistics $F_s(X, Z)$ of a sample of individuals that come from the population of interest or from populations that are somewhat related to the one of interest. This is where probability theory steps in. It allows us to assign probabilities to all the possible statistics $F(X, Z)$. From these probabilities we can calculate the *expected* value $\bar{f}(x \mid z)$ of the conditional frequencies $f(x \mid z)$. Decision theory says that the expected value $\bar{f}(x \mid z)$ should then be used, in this uncertain case, in eq. (7) in place of the unknown $f(x \mid z)$. The decision procedure (8) can then be used again.

Probability theory says that in this particular situation the probability of a particular possible statistics $F(X, Z)$ is the product of two factors having intuitive interpretations:

- the probability of observing the statistics $F_s(X, Z)$ of our data sample, assuming the full statistics to be $F(X, Z)$. With some combinatorics it can be shown that this probability is proportional to

$$\exp \left[\sum_{X,Z} F_s(X, Z) \ln F(X, Z) \right] \quad (9)$$

The argument of the exponential is the cross-entropy between $F_s(X, Z)$ and $F(X, Z)$; this is the reason of its appearance in the loss function used for classifiers⁸.

This factor tells us how much the possible statistics *fit* the sample data; it gives more weight to statistics with a better fit.

- the probability of the full statistics $F(X, Z)$ for reasons not present in the data, for example because of physical laws, biological plausibility, or similar.

This factor tells us whether the possible statistics should be favourably considered, or maybe even discarded instead, for reasons that go beyond the data we have seen; in other words, whether the hypothetical statistics would *generalize* well beyond the sample data.

The final probability comes from the balance between these ‘fit’ and ‘generalization’ factors. Note that the first factor becomes more important as the sample size and therefore $F_s(X, Z)$ increases; the sample data eventually determine what the most probable statistics is, if the sample is large enough.

A similar probabilistic reasoning applies if our sample data come not from the population of interest but from a population having at least the same *conditional* frequencies of as the one of interest, either $f(X | Z)$ or $f(Z | X)$. The latter case must be examined with care when our purpose is to guess X from Z . In this case we cannot use the conditional frequencies $f_s(X | Z)$ that appear in the data to obtain the expected value $\bar{f}(X | Z)$: they could be completely different from the ones of the population of interest. We must instead use the sample conditional frequencies $f_s(Z | X)$ to obtain the expected value $\bar{f}(Z | X)$, and then combine the latter with an appropriate probability $P(X)$ through Bayes’s theorem:

$$\frac{\bar{f}(Z | X) P(X)}{\sum_X \bar{f}(Z | X) P(X)} . \quad (10)$$

The probability $P(X)$ cannot be obtained from the data, but requires a separate study or survey. In medical applications, where X represents for example the presence or absence of a disease, the probability $P(X)$ is

⁸ Bridle 1990; MacKay 1992.

the base rate of the disease. Direct use of $f_s(X | Z)$ from the data instead of (10) is the ‘base-rate fallacy’⁹.

In supervised learning the classifier is trained to learn the most probable $f(X | Z)$ from the data. The training finds the $f(X | Z)$ that most closely fits the conditional frequency $f_s(X | Z)$ of the sampled data; this roughly corresponds to maximizing the first factor (9) described above. The architecture and the parameter regularizer of the classifier play the role of the second factor.

Bibliography

- (‘de X’ is listed under D, ‘van X’ under V, and so on, regardless of national conventions.)
- Axelsson, S. (2000): *The base-rate fallacy and the difficulty of intrusion detection*. ACM Trans. Inf. Syst. Secur. **3**³, 186–205. DOI:10.1145/357830.357849, <http://www.scs.carleton.ca/~soma/id-2007w/readings/axelsson-base-rate.pdf>.
- Baker, S. G., Pinsky, P. F. (2001): *A proposed design and analysis for comparing digital and analog mammography special receiver operating characteristic methods for cancer screening*. J. Am. Stat. Assoc. **96**⁴⁵⁴, 421–428. DOI:10.1198/016214501753168136.
- Bridle, J. S. (1990): *Probabilistic interpretation of feedforward classification network outputs, with relationships to statistical pattern recognition*. Neurocomputing **68**, 227–236. DOI: 10.1007/978-3-642-76153-9_28.
- Goodman, L. A., Kruskal, W. H. (1954): *Measures of association for cross classifications*. J. Am. Stat. Assoc. **49**²⁶⁸, 732–764. DOI:10.1080/01621459.1954.10501231. See corrections Goodman, Kruskal (1957; 1958) and also Goodman, Kruskal (1959; 1963; 1972).
- (1957): *Corrigenda: Measures of association for cross classifications*. J. Am. Stat. Assoc. **52**²⁸⁰, 578. DOI:10.1080/01621459.1957.10501415. See Goodman, Kruskal (1954).
- (1958): *Corrigenda: Measures of association for cross classifications*. J. Am. Stat. Assoc. **53**²⁸⁴, 1031. DOI:10.1080/01621459.1958.10501492. See Goodman, Kruskal (1954).
- (1959): *Measures of association for cross classifications. II: Further discussion and references*. J. Am. Stat. Assoc. **54**²⁸⁵, 123–163. DOI:10.1080/01621459.1959.10501503. See also Goodman, Kruskal (1954; 1963; 1972).
- (1963): *Measures of association for cross classifications. III: Approximate sampling theory*. J. Am. Stat. Assoc. **58**³⁰², 310–364. DOI:10.1080/01621459.1963.10500850. See correction Goodman, Kruskal (1970) and also Goodman, Kruskal (1954; 1959; 1972).
- (1970): *Corrigenda: Measures of association for cross classifications. III: Approximate sampling theory*. J. Am. Stat. Assoc. **65**³³⁰, 1011. DOI:10.1080/01621459.1970.10481142. See Goodman, Kruskal (1963).
- (1972): *Measures of association for cross classifications, IV: Simplification of asymptotic variances*. J. Am. Stat. Assoc. **67**³³⁸, 415–421. DOI:10.1080/01621459.1972.10482401. See also Goodman, Kruskal (1954; 1959; 1963).
- Hand, D., Christen, P. (2018): *A note on using the F-measure for evaluating record linkage algorithms*. Stat. Comput. **28**³, 539–547. DOI:10.1007/s11222-017-9746-6.
- Jeffrey, R. C. (1965): *The Logic of Decision*. (McGraw-Hill, New York).

⁹ Russell & Norvig 2022 § 12.5; Axelsson 2000; Jenny et al. 2018.

- Jenny, M. A., Keller, N., Gigerenzer, G. (2018): *Assessing minimal medical statistical literacy using the Quick Risk Test: a prospective observational study in Germany*. *BMJ Open* **8**, e020847, e020847corr2. DOI:10.1136/bmjopen-2017-020847, DOI:10.1136/bmjopen-2017-020847corr2.
- Kruskal, W., Mosteller, F. (1979a): *Representative sampling, I: Non-scientific literature*. *Int. Stat. Rev.* **47**¹, 13–24. See also Kruskal, Mosteller (1979b,c; 1980).
- (1979b): *Representative sampling, II: Scientific literature, excluding statistics*. *Int. Stat. Rev.* **47**², 111–127. See also Kruskal, Mosteller (1979a,c; 1980).
- (1979c): *Representative sampling, III: The current statistical literature*. *Int. Stat. Rev.* **47**³, 245–265. See also Kruskal, Mosteller (1979a,b; 1980).
- (1980): *Representative sampling, IV: The history of the concept in statistics, 1895–1939*. *Int. Stat. Rev.* **48**², 169–195. See also Kruskal, Mosteller (1979a,b,c).
- Lobo, J. M., Jiménez-Valverde, A., Real, R. (2008): *AUC: a misleading measure of the performance of predictive distribution models*. *Glob. Ecol. Biogeogr.* **17**², 145–151. DOI: 10.1111/j.1466-8238.2007.00358.x, <https://www2.unil.ch/biomapper/Download/Lobo-GloEcoBioGeo-2007.pdf>.
- MacKay, D. J. C. (1992): *The evidence framework applied to classification networks*. *Neural Comput.* **4**⁵, 720–736. <http://www.inference.phy.cam.ac.uk/mackay/PhD.html>, DOI: 10.1162/neco.1992.4.5.720.
- Matthews, B. W. (1975): *Comparison of the predicted and observed secondary structure of T4 phage lysozyme*. *Biochim. Biophys. Acta* **405**², 442–451. DOI:10.1016/0005-2795(75)90109-9.
- North, D. W. (1968): *A tutorial introduction to decision theory*. *IEEE Trans. Syst. Sci. Cybern.* **4**³, 200–210. DOI:10.1109/TSSC.1968.300114, <https://stat.duke.edu/~scs/Courses/STAT102/DecisionTheoryTutorial.pdf>.
- Provost, F. (2000): *Machine learning from imbalanced data sets 101*. Tech. rep. WS-00-05-001. (AAAI, Menlo Park, USA). <https://aaai.org/Library/Workshops/2000/ws00-05-001.php>.
- Russell, S. J., Norvig, P. (2022): *Artificial Intelligence: A Modern Approach*, Fourth Global ed. (Pearson, Harlow, UK). First publ. 1995.
- Sammut, C., Webb, G. I., eds. (2017): *Encyclopedia of Machine Learning and Data Mining*, 2nd ed. (Springer, Boston). DOI:10.1007/978-1-4899-7687-1. First publ. 2011.
- Yule, G. U. (1912): *On the methods of measuring association between two attributes*. *J. R. Stat. Soc.* **75**⁶, 579–652. DOI:10.1111/j.2397-2335.1912.tb00463.x.
- Zhu, Q. (2020): *On the performance of Matthews correlation coefficient (MCC) for imbalanced dataset*. *Pattern Recognit. Lett.* **136**, 71–80. DOI:10.1016/j.patrec.2020.03.030.